

---

# Latent Template Induction with Gumbel-CRFs

---

Yao Fu<sup>1,\*</sup>, Chuanqi Tan<sup>2</sup>, Bin Bi<sup>2</sup>, Mosha Chen<sup>2</sup>, Yansong Feng<sup>3</sup>, Alexander Rush<sup>4</sup>

<sup>1</sup>ILCC, University of Edinburgh, <sup>2</sup>Alibaba Group, <sup>3</sup>WICT, Peking University, <sup>4</sup>Cornell University  
yao.fu@ed.ac.uk, {chuanqi.tcq; b.bi; chenmosha.cms}@alibaba-inc.com  
fengyansong@pku.edu.cn, arush@cornell.edu

## Abstract

Learning to control the structure of sentences is a challenging problem in text generation. Existing work either relies on simple deterministic approaches or RL-based hard structures. We explore the use of structured variational autoencoders to infer latent templates for sentence generation using a soft, continuous relaxation in order to utilize reparameterization for training. Specifically, we propose a Gumbel-CRF, a continuous relaxation of the CRF sampling algorithm using a relaxed Forward-Filtering Backward-Sampling (FFBS) approach. As a reparameterized gradient estimator, the Gumbel-CRF gives more stable gradients than score-function based estimators. As a structured inference network, we show that it learns interpretable templates during training, which allows us to control the decoder during testing. We demonstrate the effectiveness of our methods with experiments on data-to-text generation and unsupervised paraphrase generation.

## 1 Introduction

Recent work in NLP has focused on model *interpretability* and *controllability* [64, 35, 25, 56, 17], aiming to add transparency to black-box neural networks and control model outputs with task-specific constraints. For tasks such as data-to-text generation [51, 64] or paraphrasing [38, 17], interpretability and controllability are especially important as users are interested in what linguistic properties – e.g., syntax [4], phrases [64], main entities [50] and lexical choices [17] – are controlled by the model and which part of the model controls the corresponding outputs.

Most existing work in this area relies on non-probabilistic approaches or on complex Reinforcement Learning (RL)-based hard structures. Non-probabilistic approaches include using attention weights as sources of interpretability [27, 62], or building specialized network architectures like entity modeling [50] or copy mechanism [23]. These approaches take advantages of differentiability and end-to-end training, but does not incorporate the expressiveness and flexibility of probabilistic approaches [46, 7, 31]. On the other hand, approaches using probabilistic graphical models usually involve non-differentiable sampling [30, 66, 35]. Although these structures exhibit better interpretability and controllability [35], it is challenging to train them in an end-to-end fashion.

In this work, we aim to combine the advantages of relaxed training and graphical models, focusing on conditional random field (CRF) models. Previous work in this area primarily utilizes the score function estimator (aka. REINFORCE) [63, 53, 30, 33] to obtain Monte Carlo (MC) gradient estimation for simplistic categorical models [45, 44]. However, given the combinatorial search space, these approaches suffer from high variance [21] and are notoriously difficult to train [30]. Furthermore, in a linear-chain CRF setting, score function estimators can only provide gradients for the *whole* sequence, while it would be ideal if we can derive fine-grained pathwise gradients [45] for *each step* of the sequence. In light of this, naturally one would turn to reparameterized estimators with pathwise gradients which are known to be more stable with lower variance [31, 45].

---

\*Work done during an internship at Alibaba DAMO Academy, in collaboration with PKU and Cornell.

Our simple approach for reparameterizing CRF inference is to directly relax the sampling process itself. Gumbel-Softmax [28, 39] has become a popular method for relaxing categorical sampling. We propose to utilize this method to relax each step of CRF sampling utilizing the forward-filtering backward-sampling algorithm [46]. Just as with Gumbel-Softmax, this approach becomes exact as temperature goes to zero, and provides a soft relaxation in other cases. We call this approach *Gumbel-CRF*. As is discussed by previous work that a structured latent variable may have a better inductive bias for capturing the discrete nature of sentences [29, 30, 17], we apply Gumbel-CRF as the inference model in a structured variational autoencoder for learning latent templates that control the sentence structures. Templates are defined as a sequence of states where each state controls the content (e.g., properties of the entities being discussed) of the word to be generated.

Experiments explore the properties and applications of the Gumbel-CRF approach. As a *reparameterized gradient estimator*, compared with score function based estimators, Gumbel-CRF not only gives lower-variance and fine-grained gradients for each sampling step, which leads to a better text modeling performance, but also introduce large practical advantages with significantly fewer parameters to tune and faster convergence, bring benefits to all practitioners (§ 6.1). As a *structured inference network*, like other hard models trained with REINFORCE, Gumbel-CRF also induces interpretable and controllable templates for generation. We demonstrate the interpretability and controllability on unsupervised paraphrase generation and data-to-text generation (§ 6.2). Our codes are in <https://github.com/FranxYao/Gumbel-CRF>.

## 2 Related Work

**Latent Variable Models and Controllable Text Generation.** Broadly, our model follows the line of work on deep latent variable models [15, 31, 55, 24, 12] for text generation [29, 43, 17, 68]. At an intersection of graphical models and deep learning, these works aim to embed interpretability and controllability into neural networks with continuous [8, 68], discrete [28, 39], or structured latent variables [30]. One typical template model is the Hidden Semi-Markov Model (HSMM), proposed by Wiseman et al. [64]. They use a neural generative HSMM model for joint learning the latent and the sentence with exact inference. Li and Rush [35] further equip a Semi-Markov CRF with posterior regularization [19]. While they focus on regularizing the inference network, we focus on reparameterizing it. Other works about controllability include [56, 25, 34], but many of them stay at word-level [18] while we focus on structure-level controllability.

**Monte Carlo Gradient Estimation and Continuous Relaxation of Discrete Structures.** Within the range of MC gradient estimation [45], our Gumbel-CRF is closely related to reparameterization and continuous relaxation techniques for discrete structures [65, 37, 32]. To get the MC gradient for discrete structures, many previous works use score-function estimators [53, 44, 30, 66]. This family of estimators is generally hard to train, especially for a structured model [30], while reparameterized estimators [31, 55] like Gumbel-Softmax [28, 39] give a more stable gradient estimation. In terms of continuous relaxation, the closest work is the differentiable dynamic programming proposed by Mensch and Blondel [41]. However, their approach takes an optimization perspective, and it is not straightforward to combine it with probabilistic models. Compared with their work, our Gumbel-CRF is a specific differentiable DP tailored for FFBS with Gumbel-Softmax. In terms of reparameterization, the closest work is the Perturb-and-MAP Markov Random Field (PM-MRF), proposed by Papandreou and Yuille [47]. However, when used for sampling from CRFs, PM-MRF is a biased sampler, while FFBS is unbiased. We will use a continuously relaxed PM-MRF as our baseline, and compare the gradient structures in detail in the Appendix.

## 3 Gumbel-CRF: Relaxing the FFBS Algorithm

In this section, we discuss how to relax a CRF with Gumbel-Softmax to allow for reparameterization. We start from Gumbel-Max [28], a technique for sampling from a categorical distribution. Let  $Y$  be a categorical random variable with domain  $\{1, \dots, K\}$  parameterized by the probability vector  $\pi = [\pi_1, \dots, \pi_K]$ , denoted as  $y \sim \text{Cat}(\pi)$ . Let  $G(0)$  denotes the standard Gumbel distribution,  $g_i \sim G(0), i \in \{1, \dots, K\}$  are i.i.d. gumbel noise. Gumbel-Max sampling of  $Y$  can be performed as:  $y = \arg \max_i (\log \pi_i + g_i)$ . Then the Gumbel-Softmax reparameterization is a continuous relaxation

---

**Algorithm 1** Forward Filtering Backward Sampling

---

```

1: Input:  $\Phi(z_{t-1}, z_t, x_t), t \in \{1, \dots, T\}, \alpha_{1:T}, Z$ 
2: Calculate  $p(z_T|x) = \alpha_T/Z$ 
3: Sample  $\hat{z}_T \sim p(z_T|x)$ 
4: for  $t \leftarrow T-1, 1$  do
5:    $p(z_t|\hat{z}_{t+1}, x) = \frac{\Phi(z_t, \hat{z}_{t+1}, x_{t+1})\alpha_t(z_t)}{\alpha_{t+1}(\hat{z}_{t+1})}$ 
6:   Sample  $\hat{z}_t \sim p(z_t|\hat{z}_{t+1}, x)$ 
7: end for
8: Return  $\hat{z}_{1:T}$ 

```

---




Figure 1: Gumbel-CRF FFBS algorithm and visualization for sequence-level v.s. stepwise gradients. Solid arrows show the forward pass and dashed arrows show the backward pass.

---

**Algorithm 2** Gumbel-CRF (Forward Filtering Backward Sampling with Gumbel-Softmax)

---

```

1: Input:  $\Phi(z_{t-1}, z_t, x_t), t \in \{1, \dots, T\}, \alpha_{1:T}, Z$ 
2: Calculate:
3:    $\pi_T = \alpha_T/Z$ 
4:    $\tilde{z}_T = \text{softmax}((\log \pi_T + g)/\tau), g \sim G(0)$ 
5:    $\hat{z}_T = \text{argmax}(\tilde{z}_T)$ 
6: for  $t \leftarrow T-1, 1$  do
7:    $\pi_t = \frac{\Phi(z_t, \hat{z}_{t+1}, x_{t+1})\alpha_t(z_t)}{\alpha_{t+1}(\hat{z}_{t+1})}$ 
8:    $\tilde{z}_t = \text{softmax}((\log \pi_t + g)/\tau), g \sim G(0)$ 
9:    $\hat{z}_t = \text{argmax}(\tilde{z}_t)$ 
10: end for
11: Return  $\hat{z}_{1:T}, \tilde{z}_{1:T}$   $\triangleright \tilde{z}$  is a relaxation for  $\hat{z}$ 

```

---

of Gumbel-Max by replacing the hard argmax operation with the softmax,

$$\tilde{y} = \text{softmax}([\log \pi_1 + g_1, \dots, \log \pi_K + g_K]) = \frac{\exp((\log \pi + g)/\tau)}{\sum_i \exp((\log \pi_i + g_i)/\tau)} \quad (1)$$

where  $\tilde{y}$  can be viewed as a relaxed one-hot vector of  $y$ . As  $\tau \rightarrow 0$ , we have  $\tilde{y} \rightarrow y$ .

Now we turn our focus to CRFs which generalize softmax to combinatorial structures. Given a sequence of inputs  $x = [x_1, \dots, x_T]$  a linear-chain CRF is parameterized by the factorized potential function  $\Phi(z, x)$  and defines the probability of a state sequence  $z = [z_1, \dots, z_T], z_t \in \{1, 2, \dots, K\}$  over  $x$ .

$$p(z|x) = \frac{\Phi(z, x)}{Z} \quad \Phi(z, x) = \prod_t \Phi(z_{t-1}, z_t, x_t) \quad (2)$$

$$\alpha_{1:T} = \text{Forward}(\Phi(z, x)) \quad Z = \sum_i \alpha_T(i) \quad (3)$$

The conditional probability of  $z$  is given by the Gibbs distribution with a factorized potential (equation 2). The partition function  $Z$  can be calculated with the Forward algorithm [59] (equation 3) where  $\alpha_t$  is the forward variable summarizing the potentials up to step  $t$ .

To sample from a linear-chain CRF, the standard approach is to use the forward-filtering backward-sampling (FFBS) algorithm (Algorithm 1). This algorithm takes  $\alpha$  and  $Z$  as inputs and samples  $z$  with a backward pass utilizing the locally conditional independence. The hard operation comes from the sampling operation for  $\hat{z}_t \sim p(z_t|\hat{z}_{t+1}, x)$  at each step (line 6), and this is the operation that we want to find a continuous relaxation.

Our key observation is that  $p(z_t|\hat{z}_{t+1}, x)$  is a categorical distribution, which can be directly relaxed with Gumbel-Softmax. This leads to our derivation of the Gumbelized-FFBS algorithm (Algorithm 2). The backbone of Algorithm 2 is the same as the original FFBS except for two key modifications: (1) Gumbel-Max (line 8-9) recovers the *unbiased sample*  $\hat{z}_t$  and the same sampling path as Algorithm 1; (2) the continuous relaxation of argmax with softmax (line 8) that gives the differentiable<sup>2</sup> (but biased)  $\tilde{z}$ .

We can apply this approach in any setting requiring structured sampling. For instance let  $p_\phi(z|x)$  denote the sampled distribution and  $f(z)$  be a downstream model. We achieve a reparameterized gradient estimator for the sampled model with  $\tilde{z}$ :

$$\nabla_\phi \mathbb{E}_{p_\phi(z|x)}[f(z)] \approx \mathbb{E}_{g \sim G(0)}[\nabla_\phi f(\tilde{z}(\phi, g))] \quad (4)$$

---

<sup>2</sup>Note that argmax is also differentiable almost everywhere, however its gradient is almost 0 everywhere and not well-defined at the jumping point [49]. Our relaxed  $\tilde{z}$  does not have these problems.

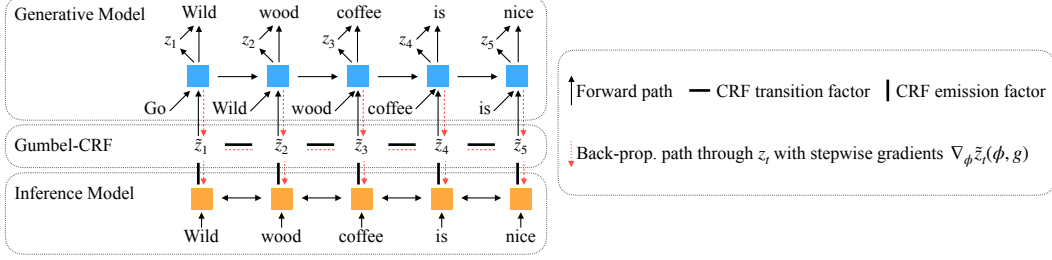


Figure 2: Architecture of our model. Note the structure of gradients induced by Gumbel-CRF differs significantly with REINFORCE: REINFORCE induces gradient for the whole sequence  $\nabla_{\phi} \log q_{\phi}(z|x)$  while Gumbel-CRF allows the model to back-prop gradients along each sample steps  $\nabla_{\phi} \tilde{z}_t$  (red dashed arrows).

We further consider a straight-through (ST) version [28] of this estimator where we use the hard sample  $\hat{z}$  in the forward pass, and back-propagate through the soft  $\tilde{z}$ .

We highlight that one advantage of this reparameterized estimator (and its ST version), compared with the score-function estimator, is that  $\tilde{z}$  induces fine-grained gradients for each step  $\tilde{z}_t$ . As illustrated in Figure 1,  $\tilde{z}_t$  is a differentiable function of the potential and the noise:  $\tilde{z}_t = \tilde{z}_t(\phi, g)$ . So during back-propagation, we can take *stepwise gradients*  $\nabla_{\phi} \tilde{z}_t(\phi, g)$ . On the other hand, with a score-function estimator, we only observe the reward for the whole sequence, so the gradient is at the *sequence level*  $\nabla_{\phi} \log p_{\phi}(z|x)$ . The lack of intermediate reward, i.e., stepwise gradients, is an essential challenge in reinforcement learning [57, 67]. This challenge is circumvented with the structure of Gumbel-CRF, thus significantly reducing modeling complexity in practice (detailed demonstrations in experiments).

#### 4 Gumbel-CRF VAE

An appealing use of reparameterizable CRF models is to learn variational autoencoders (VAEs) with a structured inference network. Past work has shown that these models (trained with RL) [30, 35] can learn to induce useful latent structure while producing accurate models. We introduce a VAE for learning latent templates for text generation, starting with the generative model in an unconditional setting.

**Generative Model** We assume a simple generative process where each words  $x_t$  of a sentence  $x = [x_1, \dots, x_T]$  is controlled by a sequence of latent states  $z = [z_1, \dots, z_T]$ , i.e., template, similar to Li and Rush [35]:

$$p_{\theta}(x, z) = \prod_t p(x_t | z_t, z_{1:t-1}, x_{1:t-1}) \cdot p(z_t | z_{1:t-1}, x_{1:t-1}) \quad (5)$$

$$h_t = \text{Dec}([z_{t-1}; x_{t-1}], h_{t-1}) \quad (6)$$

$$p(z_t | z_{1:t-1}, x_{1:t-1}) = \text{Softmax}(\text{FF}(h_t)) \quad (7)$$

$$p(x_t | z_t, z_{1:t-1}, x_{1:t-1}) = \text{Softmax}(\text{FF}([e(z_t); h_t])) \quad (8)$$

Where  $\text{Dec}(\cdot)$  denotes the decoder,  $\text{FF}(\cdot)$  denotes a feed-forward network,  $h_t$  denotes the decoder state,  $[\cdot; \cdot]$  denotes vector concatenation.  $e(\cdot)$  denotes the embedding function. Under this formulation, the generative model is autoregressive w.r.t. both  $x$  and  $z$ . Intuitively, it generates the control states and words in turn, and the current word  $x_t$  is primarily controlled by its corresponding  $z_t$ .

**Inference Model** Since the exact inference of  $p_{\theta}(z|x)$  is intractable, we approximate it with a structured variational family and optimize the following form of the ELBO objective:

$$\text{ELBO} = \mathbb{E}_{q_{\phi}(z|x)}[\log p_{\theta}(x, z)] - \mathcal{H}[q_{\phi}(z|x)] \quad (9)$$

Where  $\mathcal{H}$  denotes the entropy. The key use of Gumbel-CRF is for the inference model  $q_{\phi}(z|x)$  to learn templates. Following past work [26], we parameterize it as a CRF whose potential is predicted by a neural encoder:

$$h_{1:T}^{(\text{enc})} = \text{Enc}(x_{1:T}) \quad \Phi(z_t, x_t) = W_{\Phi} h_t^{(\text{enc})} + b_{\Phi} \quad (10)$$

$$\Phi(z_{t-1}, z_t, x_t) = \Phi(z_{t-1}, z_t) \cdot \Phi(z_t, x_t) \quad (11)$$

Where  $\text{Enc}(\cdot)$  denotes the encoder and  $h_t^{(\text{enc})}$  is the encoder state. With this formulation, the entropy term in equation 9 can be computed efficiently with dynamic programming, which is differentiable by itself [40].

**Extension to Conditional Settings** For practical applications like paraphrasing and data-to-text, we make a conditional extension where the generative model is conditioned on a source data structure  $s$ , formulated as  $p_\theta(x, z|s)$ . Specifically, for paraphrase generation,  $s = [s_1, \dots, s_N]$  is the bag of words (a set,  $N$  being the size of the BOW) of the source sentence, similar to Fu et al. [17]. We aim to generate a different sentence  $x$  with the same meaning as the input sentence. In addition to being autoregressive on  $x$  and  $z$ , the decoder also attend to [1] and copy [23] from  $s$ . For data-to-text, we denote the source data is formed as a table of key-value pairs:  $s = [(k_1, v_1), \dots, (k_N, v_N)]$ ,  $N$  being size of the table. We aim to generate a sentence  $x$  that best describes the table. Again, we condition the generative model on  $s$  by attending to and copying from it. Note our formulation would effectively become *a neural version of slot-filling*: for paraphrase generation we fill the BOW into the neural templates, and for data-to-text we fill the values into neural templates. We assume the inference model is independent from the source  $s$  and keep it unchanged, i.e.,  $q_\phi(z|x, s) = q_\phi(z|x)$ . The ELBO objective in this conditional setting is:

$$\text{ELBO} = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x, z|s)] - \mathcal{H}[q_\phi(z|x)] \quad (12)$$

**Training and Testing** During training, for the Gumbel-CRF gradient estimator, we use the relaxed sample  $\tilde{z}$  from the Gumbelized FFBS (Algorithm 2) in both the forward and backward passes. For the ST version of Gumbel-CRF, we use the exact sample  $\hat{z}$  in the forward pass and back-propagate gradients through the relaxed  $\tilde{z}$ . At the beginning of training, to prevent the decoder from ignoring  $z$ , we apply word dropout [8], i.e., to randomly set the input word embedding at certain steps to be 0. After  $z$  converges to a meaningful local optimal, we gradually decrease word dropout ratio to 0 and recover the full model. For optimization, we add a  $\beta$  coefficient to the entropy term, as is in the  $\beta$ -VAE [24]. As is in many VAE works [8, 11], we observe the posterior collapse problem where  $q(z|x)$  converges to meaningless local optimal. We observe two types of collapsed posterior in our case: a constant posterior ( $q$  outputs a fixed  $z$  no matter what  $x$  is. This happens when  $\beta$  is too weak), and a uniform posterior (when  $\beta$  is too strong). To prevent posterior collapse,  $\beta$  should be carefully tuned to achieve a balance. During testing, for evaluating paraphrasing and data-to-text, we use greedy decoding for both  $z$  and  $x$ ; for controlling the structure of generated sentences, we sample a fixed MAP  $z$  from the training set (i.e., the aggregated variational posterior), feed it to each decoder steps, and use it to control the generated  $x$ .

## 5 Experimental Setup

Our experiments are in two steps. First, as a reparameterized gradient estimator, we compare Gumbel-CRF to other common gradient estimators on the standard text modeling task, and show it gives a better density estimation performance with lower variance. Then we integrate Gumbel-CRF to real-world models, specifically paraphrase generation and data-to-text generation. We show that like the hard models trained with REINFORCE [35], models trained with Gumbel-CRF also induce meaningful templates for interpretability and controllability with similar or better end-task performance (e.g., BLEU). Furthermore, we show that the simplicity of Gumbel-CRF gives many practical benefits including fewer parameters to tune and faster training time.

**Datasets** We focus on two datasets. For text modeling and data-to-text generation, we use the E2E dataset[51], a common dataset for learning structured templates for text [64, 35]. This dataset contains approximately 42K training, 4.6K validation and 4.6K testing sentences. The vocabulary size is 945. For paraphrase generation we follow the same setting as Fu et al. [17], and use the common MSCOCO dataset. This dataset has 94K training and 23K testing instances. The vocabulary size is 8K.


**Metrics** For evaluating the gradient estimator performance, we follow the common practice and primarily compare the test negative log-likelihood (NLL) estimated with importance sampling. We also report relative metrics: ELBO, perplexity (PPL), and entropy of the inference network. Importantly, to make all estimates unbiased, all models are evaluated in *a discrete setting with unbiased hard samples*. For paraphrase task performance, we follow Liu et al. [38], Fu et al. [17] and use BLEU (bigram to 4-gram) [48] and ROUGE [36] (R1, R2 and RL) to measure the generation quality. We note that although being widely used, the two metrics do not penalize the similarity

Table 1: Density Estimation Results. NLL is estimated with 100 importance samples. Models are selected from 3 different random seeds based on validation NLL. All metrics are evaluated on the discrete (exact) model.

| Model                | Neg. ELBO | NLL   | PPL   | Ent. | #sample |
|----------------------|-----------|-------|-------|------|---------|
| RNNLM                | -         | 34.69 | 4.94  | -    | -       |
| PM-MRF               | 69.15     | 50.22 | 10.41 | 4.11 | 1       |
| PM-MRF-ST            | 53.16     | 37.03 | 5.48  | 2.04 | 1       |
| REINFORCE-MS         | 35.11     | 34.50 | 4.84  | 3.48 | 5       |
| REINFORCE-MS-C       | 34.35     | 33.82 | 4.71  | 3.34 | 5       |
| Gumbel-CRF (ours)    | 38.00     | 35.41 | 4.71  | 3.03 | 1       |
| Gumbel-CRF-ST (ours) | 34.18     | 33.13 | 4.54  | 3.26 | 1       |

| Estimators     | Score /Reparam. | Seq. Level/ Stepwise | Unbiased MC Sample | Unbiased Grad. |
|----------------|-----------------|----------------------|--------------------|----------------|
| REINFORCE-MS   | Score           | Seq.                 | Unbiased           | Unbiased       |
| REINFORCE-MS-C | Score           | Seq.                 | Unbiased           | Unbiased       |
| PM-MRF         | Reparam.        | Step                 | Biased             | Biased         |
| PM-MRF-ST      | Reparam.        | Step                 | Biased             | Biased         |
| Gumbel-CRF     | Reparam.        | Step                 | Biased             | Biased         |
| Gumbel-CRF-ST  | Reparam.        | Step                 | Unbiased           | Biased         |

(A) Characteristics of the estimators we compare



(B) Log variance ratio, training curve

Figure 3: Text Modeling. (A). Characteristics of gradient estimators. (B). Variance comparison, Gumbel-CRF vs REINFORCE-MS, training curve.

between the generated sentence and the input sentence (because we do not want the model to simply copy the input). So we adopt iBLUE [58], a specialized BLUE score that penalize the ngram overlap between the generated sentence and the input sentence, and use it as our primary metrics. The iBLUE score is defined as:  $iB(i, o, r) = \alpha B(o, r) + (1 - \alpha)B(o, i)$ , where  $iB(\cdot)$  denotes iBLUE score,  $B(\cdot)$  denotes BLUE score,  $i, o, r$  denote input, output, and reference sentences respectively. We follow Liu et al. [38] and set  $\alpha = 0.9$ . For text generation performance, we follow Li and Rush [35] and use the E2E official evaluation script, which measures BLEU, NIST [5], ROUGE, CIDEr [61], and METEOR [3]. More experimental details are in the Appendix.

## 6 Results

### 6.1 Gumbel-CRF as Gradient Estimator: Text Modeling

We compare our Gumbel-CRF (original and ST variant) with two sets of gradient estimators: score function based and reparameterized. For score function estimators, we compare our model with REINFORCE using the mean reward of other samples (MS) as baseline, as is used in Kim et al. [30], Li and Rush [35]. We further find that adding a carefully tuned constant baseline helps with the scale of the gradient (REINFORCE MS-C). For reparameterized estimators, we use a tailored Perturb-and-Map Markov Random Field (PM-MRF) estimator [47] with the continuous relaxation introduced in Corro and Titov [10]. Compared to our Gumbel-CRF, PM-MRF adds Gumbel noise to local potentials, then runs a relaxed structured argmax algorithm [41]. We further consider a straight-through (ST) version of PM-MRF. The basics of these estimators can be characterized from four dimensions, as listed in Figure 3(A). We encourage the readers to go through the appendix for an in-depth theoretical comparison of gradient structures between these estimators.

Table 1 shows the main results comparing different gradient estimators on text modeling. Our Gumbel-CRF-ST outperforms other estimators in terms of NLL and PPL. With fewer samples required, reparameterization and continuous relaxation used in Gumbel-CRF are particularly effective for learning structured inference networks. We also see that PM-MRF estimators perform worse than other estimators. Due to the biased nature of the Perturb-and-MAP sampler, during optimization, PM-MRF is *not optimizing the actual model*. As a sampler (rather than a gradient estimator) Gumbel-

Table 2: Paraphrase Generation. Upper: supervised models, Lower: unsupervised models. Models are selected from 5 random seeds based validation iB4 (iBLUE 4 gram).

| Model                      | iB4   | B2    | B3    | B4    | R1    | R2    | RL    |
|----------------------------|-------|-------|-------|-------|-------|-------|-------|
| LBOW [17]                  | -     | 51.14 | 35.66 | 25.27 | 42.08 | 16.13 | 38.16 |
| Gaussian VAE[8]            | 7.48  | 24.90 | 13.04 | 7.29  | 22.05 | 4.64  | 26.05 |
| CGMH [42]                  | 7.84  | -     | -     | 11.45 | 32.19 | 8.67  | -     |
| UPSA [38]                  | 9.26  | -     | -     | 14.16 | 37.18 | 11.21 | -     |
| Ours trained w. REINFORCE  | 11.20 | 41.29 | 26.54 | 17.10 | 32.57 | 10.20 | 34.97 |
| Ours trained w. Gumbel-CRF | 10.20 | 38.98 | 24.65 | 15.75 | 31.10 | 9.24  | 33.60 |

Table 3: Data-to-text generation results. Upper: neural models, Lower: template-related models. Models are selected from 5 different random seeds based on validation BLEU.

| Model                      | BLEU  | NIST | ROUGE | CIDEr | METEOR |
|----------------------------|-------|------|-------|-------|--------|
| D&J[14]                    | 65.93 | 8.59 | 68.50 | 2.23  | 44.83  |
| KV2Seq[16]                 | 74.72 | 9.30 | 70.69 | 2.23  | 46.15  |
| SUB[14]                    | 43.78 | 6.88 | 54.64 | 1.39  | 37.35  |
| HMM[64]                    | 55.17 | 7.14 | 65.70 | 1.70  | 41.91  |
| HMM-AR[64]                 | 59.80 | 7.56 | 65.01 | 1.95  | 38.75  |
| SM-CRF PC [35]             | 67.12 | 8.52 | 68.70 | 2.24  | 45.40  |
| Ours trained w. REINFORCE  | 60.41 | 7.99 | 62.54 | 1.78  | 38.04  |
| Ours trained w. Gumbel-CRF | 65.83 | 8.43 | 65.06 | 1.98  | 41.44  |

CRF is less biased than PM-MRF. We further observe that both the ST version of Gumbel-CRF and PM-MRF perform better than the non-ST version. We posit that this is because of the consistency of using hard samples in both training and testing time (although non-ST has other advantages).

**Variance Analysis** To show that reparameterized estimators have lower variance, we compare the log variance ratio of Gumbel-CRF and REINFORCE-MS-C (Figure 3 B), which is defined as  $r = \log(\text{Var}(\nabla_{\phi}\mathcal{L})/|\nabla_{\phi}\mathcal{L}|)$  ( $\nabla_{\phi}\mathcal{L}$  is gradients of the inference model)<sup>3</sup>. We see that Gumbel-CRF has a lower training curve of  $r$  than REINFORCE-MS, showing that it is more stable for training.

## 6.2 Gumbel-CRF for Control: Paraphrase Generation and Data-to-Text

**Unsupervised Paraphrase Generation** Unsupervised paraphrase generation is defined as generating different sentences conveying the same meaning of an input sentence without parallel training instances. To show the effectiveness of Gumbel-CRF as a gradient estimator, we compare the results when our model is trained with REINFORCE. To show the overall performance of our structured model, we compare it with other unsupervised models, including: a Gaussian VAE for paraphrasing [8]; CGMH [42], a general-purpose MCMC method for controllable generation; UPSA [38], a strong paraphrasing model with simulated annealing. To better position our template model, we also report the supervised performance of a state-of-the-art latent bag of words model (LBOW) [17].

Table 2 shows the results. As expected, the supervised model LBOW performs better than all unsupervised models. Among unsupervised models, the best iB4 results come from our model trained with REINFORCE. When trained with Gumbel-CRF, our model performs slightly worse than REINFORCE, but still outperforms other paraphrasing models. We note that it is quite interesting that our model is conceptually very simple: it just view the input sentence as a bag of words, and (re-)organize them into a new sentence. Yet this simple approach outperforms models involving complicated techniques like UPSA on certain metrics like iBLUE. This observation may open the question about what kind of inductive bias is actually required for unsupervised paraphrasing. To summarize, like REINFORCE, our Gumbel-CRF is also effective for training paraphrasing models

<sup>3</sup>Previous works compare variance, rather than variance ratio [60, 20]. We think that simply comparing the variance is only reasonable when the scale of gradients are approximately the same, which may not hold in different estimators. In our experiments, we observe that the gradient scale of Gumbel-CRF is significantly smaller than REINFORCE, thus the variance ratio may be a better proxy for measuring stability.

Table 4: Practical benefits of using Gumbel-CRF. Typically, REINFORCE has a long list of parameters to tune:  $h$  entropy regularization,  $b_0$  constant baseline,  $b$  baseline model,  $r$  reward scaling,  $\#s$  number of MC sample. Gumbel-CRF reduces the engineering complexity with significantly less parameters to tune ( $h$  entropy regularization,  $\tau$  temperature annealing), less sample required (thus less memory consumption), and less time consumption. Models tested on Nvidia P100 with batch size 100.

| Model      | Hyperparams.        | #s | GPU mem | Sec. per batch |
|------------|---------------------|----|---------|----------------|
| REINFORCE  | $h, b_0, b, r, \#s$ | 5  | 1.8G    | 1.42           |
| Gumbel-CRF | $h, \tau$           | 1  | 1.1G    | 0.48           |

|  |
|--|
| name: clowns   eatype: coffee shop   food: chinese   customer_rating: 1 out of 5   area: riverside   near: clare hall  |
| 1. [there is a] <sub>20</sub> [coffee shop] <sub>35</sub> [in the] <sub>9</sub> [riverside] <sub>35</sub> [area .] <sub>12</sub> [serves] <sub>20</sub> [chinese] <sub>35</sub> [food] <sub>12</sub> [ . it is] <sub>20</sub> [called] <sub>35</sub> [clowns] <sub>44</sub> [ . is] <sub>20</sub> [near] <sub>35</sub> [clare hall] <sub>44</sub> [ . It has a customer rating] <sub>30</sub> [of 1 out of 5] <sub>8</sub> [ .] <sub>20</sub>                            |
| 2. [clowns] <sub>44</sub> [is a] <sub>20</sub> [expensive] <sub>12</sub> [coffee shop] <sub>35</sub> [located] <sub>12</sub> [in] <sub>9</sub> [riverside] <sub>35</sub> [area] <sub>12</sub> [ .] <sub>20</sub>   |
| 3. [clowns] <sub>44</sub> [is a] <sub>20</sub> [coffee shop] <sub>35</sub> [in the riverside] <sub>35</sub> [ . it is] <sub>20</sub> [family friendly] <sub>12</sub> [and has a] <sub>20</sub> [1] <sub>45</sub> [out of 5] <sub>8</sub> [stars] <sub>12</sub> [rating .] <sub>20</sub>  |
| name: browns cambridge   eatype: coffee shop   food: chinese   customer_rating: 1 out of 5   area: riverside   familyfriendly: yes   near: crowne plaza hotel  |
| 1. [browns cambridge] <sub>44</sub> [offers] <sub>12</sub> [chinese] <sub>35</sub> [food] <sub>12</sub> [near] <sub>35</sub> [crowne plaza hotel] <sub>44</sub> [in] <sub>35</sub> [riverside] <sub>35</sub> [ . it is a] <sub>20</sub> [coffee shop] <sub>35</sub> [ . not children friendly] <sub>12</sub> [and has a] <sub>20</sub> [5] <sub>45</sub> [out of 5] <sub>8</sub> [rating .] <sub>20</sub>  |
| 2. [there is a] <sub>20</sub> [moderately priced restaurant] <sub>12</sub> [that serves] <sub>20</sub> [chinese] <sub>35</sub> [food] <sub>12</sub> [called] <sub>35</sub> [browns cambridge] <sub>44</sub> [coffee] <sub>35</sub> [ . it has a customer rating] <sub>30</sub> [of 5 out of 5] <sub>8</sub> [it is] <sub>20</sub> [not family-friendly] <sub>12</sub> [ . it is] <sub>20</sub> [located] <sub>12</sub> [near] <sub>35</sub> [crowne plaza] <sub>44</sub> |
| 3. [browns cambridge] <sub>44</sub> [is a] <sub>20</sub> [chinese coffee shop] <sub>35</sub> [located] <sub>12</sub> [in] <sub>9</sub> [riverside near] <sub>35</sub> [crowne plaza hotel] <sub>44</sub> [ . it has a] <sub>20</sub> [customer rating] <sub>30</sub> [of 5 out of] <sub>8</sub> [5] <sub>44</sub> [and is] <sub>20</sub> [not family-friendly] <sub>12</sub> [ .] <sub>20</sub>  |

Figure 4: Controllable generation with templates.

| Bigram | Sentence Segments                    | 4gram       | Sentence Segments                    |                      |
|--------|--------------------------------------|-------------|--------------------------------------|----------------------|
| 20-12  | is located                           | 12-20-12-20 | food with a price range of           | Adjective Phrase     |
|        | and is located                       |             | food at a price range of             | Prepositional Phrase |
|        | is a family friendly                 |             | price range and is family friendly   |                      |
| 12-35  | located near                         | 35-44-12-20 | near the city centre                 | Noun Phrase          |
|        | restaurant near                      |             | near café rouge, there is a          |                      |
|        | restaurant located near              |             | in the city centre. it is            | Verbal Phrase        |
| 20-8   | has a customer rating of             | 44-20-35-20 | french food at a moderate            |                      |
|        | it has a customer rating of 5 out of |             | indian food for a moderate           |                      |
|        | and with a customer rating of        |             | fast food restaurant with a moderate |                      |

Figure 5: Template Interpretability. State ngrams capture phrasal structures.

without much lose of end-performance. We will further highlight the additional practical benefits of Gumbel-CRF in later sections.

**Data-to-Text Generation** Data-to-text generation is defined as to generate descriptions for given tabular information. Classical approaches use rule-based templates with better interpretability, while recent approaches use neural models for better performance. Here we aim to study the interpretability and controllability of latent templates. We compare our model with neural and template-based models. Neural models include: D&J[14] (with basic seq2seq); and KV2Seq[16] (with SOTA neural memory architectures); Template models include: SUB[14] (with rule-based templates); HSMM[64] (with neural templates); and SM-CRF-PC[35] (with neural templates and posterior regularization[19]).

Table 3 shows the results of data-to-text generation. As expected, neural models like KV2Seq with advanced architectures achieve the best performance. Template-related models all come with a certain level of performance costs for better controllability. Among the template-related models, SM-CRF PC performs best. However, it utilizes multiple weak supervision to achieve better template-data alignment, while our model is fully unsupervised for the template. Our model, either trained with REINFORCE or Gumbel-CRF, outperforms the baseline HSMM model. We further see that in this case, models trained with Gumbel-CRF gives better end performance than REINFORCE.

To see how the learned templates induce controllability, we conduct a qualitative study. To use the templates, after convergence, we collect and store all the MAP  $z$  for the training sentences. During testing, given an input table, we retrieve a template  $z$  and use this  $z$  as the control state for the decoder. Figure 4 shows sentences generated from templates. We can see that sentences with different



templates exhibit different structures. E.g., the first sentence for the *clowns* coffee shop starts with the location, while the second starts with the price. We also observe a *state-word correlation*. E.g., state 44 always corresponds to the name of a restaurant and state 8 always corresponds to the rating.

We further interpret the templates by associating frequent state ngrams with their corresponding sentence segments (Figure 5). like a topic model [6], semantically similar segments are clustered to the same ngrams, revealing different phrasal structures. Combining this observation with the previous state-word correlation, we can say that a template jointly encodes the lexical and structural information, where the lexical information is encoded by individual states, and the structural information is encoded by the sequential combination (ngrams) of states.

**Practical Benefits** Although our model can be trained on either REINFORCE or Gumbel-CRF, we emphasize that training structured variables with REINFORCE is notoriously difficult [35], and Gumbel-CRF substantially reduces the complexity. Table 4 demonstrates this empirically. Gumbel-CRF requires fewer hyperparameters to tune, fewer MC samples, less GPU memory, and faster training. These advantages would considerably benefit all practitioners with significantly less training time and resource consumption.

## 7 Conclusion

In this work, we propose a pathwise gradient estimator for sampling from CRFs which exhibits lower variance and more stable training than existing baselines. We apply this gradient estimator to the task of text modeling, where we use a structured inference network based on CRFs to learn latent templates. Just as REINFORCE, models trained with Gumbel-CRF can also learn meaningful latent templates that successfully encode lexical and structural information of sentences, thus inducing interpretability and controllability for text generation. Furthermore, the Gumbel-CRF gives significant practical benefits than REINFORCE, making it more applicable to real-world tasks.

## Broader Impact

Generally, this work is about controllable text generation. When applying this work to chatbots, one may get a better generation quality. This could potentially improve the accessibility [9, 54] for people who need a voice assistant to use an electronic device, e.g. people with visual, intellectual, and other disabilities [52, 2]. However, if not properly tuned, this model may generate improper sentences like fake information, putting the user at a disadvantage. Like many other text generation models, if trained with improper data (fake news [22], words of hatred [13]), a model could generate these sentences as well. In fact, one of the motivations for controllable generation is to avoid these situations [64, 13]. But still, researchers and engineers need to be more careful when facing these challenges.

## References

- [1] Dzmitry Bahdanau, Kyunghyun Cho, and Yoshua Bengio. Neural machine translation by jointly learning to align and translate. In *3rd International Conference on Learning Representations, ICLR 2015*, 2015.
- [2] Saminda Sundeepa Balasuriya, Laurianne Sitbon, Andrew A. Bayor, Maria Hoogstrate, and Margot Brereton. Use of voice activated interfaces by people with intellectual disability. *Proceedings of the 30th Australian Conference on Computer-Human Interaction*, 2018.
- [3] Satanjeev Banerjee and Alon Lavie. Meteor: An automatic metric for mt evaluation with improved correlation with human judgments. In *Proceedings of the acl workshop on intrinsic and extrinsic evaluation measures for machine translation and/or summarization*, pages 65–72, 2005.
- [4] Yu Bao, Hao Zhou, Shujian Huang, Lei Li, Lili Mou, Olga Vechtomova, Xin-yu Dai, and Jiajun Chen. Generating sentences from disentangled syntactic and semantic spaces. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics*,

- pages 6008–6019, Florence, Italy, July 2019. Association for Computational Linguistics. doi: 10.18653/v1/P19-1602. URL <https://www.aclweb.org/anthology/P19-1602>.
- [5] Anja Belz and Ehud Reiter. Comparing automatic and human evaluation of nlg systems. In *11th Conference of the European Chapter of the Association for Computational Linguistics*, 2006.
  - [6] David M Blei, Andrew Y Ng, and Michael I Jordan. Latent dirichlet allocation. *Journal of machine Learning research*, 3(Jan):993–1022, 2003.
  - [7] David M. Blei, Alp Kucukelbir, and Jon D. McAuliffe. Variational inference: A review for statisticians. *ArXiv*, abs/1601.00670, 2016.
  - [8] Samuel R Bowman, Luke Vilnis, Oriol Vinyals, Andrew M Dai, Rafal Jozefowicz, and Samy Bengio. Generating sentences from a continuous space. *arXiv preprint arXiv:1511.06349*, 2015.
  - [9] Eric Corbett and Astrid Weber. What can i say?: addressing user experience challenges of a mobile voice user interface for accessibility. *Proceedings of the 18th International Conference on Human-Computer Interaction with Mobile Devices and Services*, 2016.
  - [10] Caio Corro and Ivan Titov. Differentiable perturb-and-parse: Semi-supervised parsing with a structured variational autoencoder. *arXiv preprint arXiv:1807.09875*, 2018.
  - [11] Adji B Dieng, Yoon Kim, Alexander M Rush, and David M Blei. Avoiding latent variable collapse with generative skip models. *arXiv preprint arXiv:1807.04863*, 2018.
  - [12] Carl Doersch. Tutorial on variational autoencoders. *arXiv preprint arXiv:1606.05908*, 2016.
  - [13] Cícero Nogueira dos Santos, Igor Melnyk, and Inkit Padhi. Fighting offensive language on social media with unsupervised text style transfer. In *ACL*, 2018.
  - [14] Ondřej Dušek and Filip Jurčiček. Sequence-to-sequence generation for spoken dialogue via deep syntax trees and strings. *arXiv preprint arXiv:1606.05491*, 2016.
  - [15] Yao Fu. Deep generative models for natural language processing. URL <https://github.com/FranxYao/Deep-Generative-Models-for-Natural-Language-Processing>.
  - [16] Yao Fu and Yansong Feng. Natural answer generation with heterogeneous memory. In *Proceedings of the 2018 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long Papers)*, pages 185–195, 2018.
  - [17] Yao Fu, Yansong Feng, and John P. Cunningham. Paraphrase generation with latent bag of words. In *NeurIPS*, 2019.
  - [18] Yao Fu, Hao Zhou, Jiaze Chen, and Lei Li. Rethinking text attribute transfer: A lexical analysis. *arXiv preprint arXiv:1909.12335*, 2019.
  - [19] Kuzman Ganchev, Joao Graca, Jennifer Gillenwater, and Ben Taskar. Posterior regularization for structured latent variable models. *Journal of Machine Learning Research*, 11(67):2001–2049, 2010. URL <http://jmlr.org/papers/v11/ganchev10a.html>.
  - [20] Will Grathwohl, Dami Choi, Yuhuai Wu, Geoffrey Roeder, and David Duvenaud. Backpropagation through the void: Optimizing control variates for black-box gradient estimation. *ArXiv*, abs/1711.00123, 2018.
  - [21] Evan Greensmith, Peter L Bartlett, and Jonathan Baxter. Variance reduction techniques for gradient estimates in reinforcement learning. *Journal of Machine Learning Research*, 5(Nov): 1471–1530, 2004.
  - [22] Nir Grinberg, Kenneth Joseph, Lisa Friedland, Briony Swire-Thompson, and David Lazer. Fake news on twitter during the 2016 u.s. presidential election. *Science*, 363:374–378, 2019.
  - [23] Jiatao Gu, Zhengdong Lu, Hang Li, and Victor OK Li. Incorporating copying mechanism in sequence-to-sequence learning. *arXiv preprint arXiv:1603.06393*, 2016.

- [24] Irina Higgins, Loïc Matthey, Arka Pal, Christopher Burgess, Xavier Glorot, Matthew Botvinick, Shakir Mohamed, and Alexander Lerchner. beta-vae: Learning basic visual concepts with a constrained variational framework. In *ICLR*, 2017.
- [25] Zhiting Hu, Zichao Yang, Xiaodan Liang, Ruslan Salakhutdinov, and Eric P Xing. Toward controlled generation of text. In *Proceedings of the 34th International Conference on Machine Learning-Volume 70*, pages 1587–1596. JMLR. org, 2017.
- [26] Zhiheng Huang, Wei Xu, and Kai Yu. Bidirectional lstm-crf models for sequence tagging. *arXiv preprint arXiv:1508.01991*, 2015.
- [27] Sarthak Jain and Byron C. Wallace. Attention is not explanation. *ArXiv*, abs/1902.10186, 2019.
- [28] Eric Jang, Shixiang Gu, and Ben Poole. Categorical reparameterization with gumbel-softmax. *arXiv preprint arXiv:1611.01144*, 2016.
- [29] Yoon Kim, Sam Wiseman, and Alexander M Rush. A tutorial on deep latent variable models of natural language. *arXiv preprint arXiv:1812.06834*, 2018.
- [30] Yoon Kim, Alexander M Rush, Lei Yu, Adhiguna Kuncoro, Chris Dyer, and Gábor Melis. Unsupervised recurrent neural network grammars. *arXiv preprint arXiv:1904.03746*, 2019.
- [31] Diederik P Kingma and Max Welling. Auto-encoding variational bayes. *arXiv preprint arXiv:1312.6114*, 2013.
- [32] Wouter Kool, Herke Van Hoof, and Max Welling. Stochastic beams and where to find them: The gumbel-top-k trick for sampling sequences without replacement. *arXiv preprint arXiv:1903.06059*, 2019.
- [33] Bowen Li, Jianpeng Cheng, Yang Liu, and Frank Keller. Dependency grammar induction with a neural variational transition-based parser. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 33, pages 6658–6665, 2019.
- [34] Juncen Li, Robin Jia, He He, and Percy Liang. Delete, retrieve, generate: A simple approach to sentiment and style transfer. *arXiv preprint arXiv:1804.06437*, 2018.
- [35] Xiang Lisa Li and Alexander M. Rush. Posterior control of blackbox generation. *ArXiv*, abs/2005.04560, 2020.
- [36] Chin-Yew Lin and Eduard Hovy. Manual and automatic evaluation of summaries. In *Proceedings of the ACL-02 Workshop on Automatic Summarization-Volume 4*, pages 45–51. Association for Computational Linguistics, 2002.
- [37] Scott W Linderman, Gonzalo E Mena, Hal Cooper, Liam Paninski, and John P Cunningham. Reparameterizing the birkhoff polytope for variational permutation inference. *arXiv preprint arXiv:1710.09508*, 2017.
- [38] Xianggen Liu, Lili Mou, Fandong Meng, Hao Zhou, Jie Zhou, and Sen Song. Unsupervised paraphrasing by simulated annealing. *ArXiv*, abs/1909.03588, 2019.
- [39] Chris J Maddison, Andriy Mnih, and Yee Whye Teh. The concrete distribution: A continuous relaxation of discrete random variables. *arXiv preprint arXiv:1611.00712*, 2016.
- [40] Gideon S Mann and Andrew McCallum. Efficient computation of entropy gradient for semi-supervised conditional random fields. Technical report, MASSACHUSETTS UNIV AMHERST DEPT OF COMPUTER SCIENCE, 2007.
- [41] Arthur Mensch and Mathieu Blondel. Differentiable dynamic programming for structured prediction and attention. *arXiv preprint arXiv:1802.03676*, 2018.
- [42] Ning Miao, Hao Zhou, Lili Mou, Rui Yan, and Lei Li. Cgmh: Constrained sentence generation by metropolis-hastings sampling. In *AAAI*, 2018.
- [43] Yishu Miao. *Deep generative models for natural language processing*. PhD thesis, University of Oxford, 2017.

- [44] Andriy Mnih and Danilo Jimenez Rezende. Variational inference for monte carlo objectives. In *ICML*, 2016.
- [45] Shakir Mohamed, Mihaela Rosca, Michael Figurnov, and Andriy Mnih. Monte carlo gradient estimation in machine learning. *ArXiv*, abs/1906.10652, 2019.
- [46] Kevin P Murphy. *Machine learning: a probabilistic perspective*. 2012.
- [47] George Papandreou and Alan L Yuille. Perturb-and-map random fields: Using discrete optimization to learn and sample from energy models. In *2011 International Conference on Computer Vision*, pages 193–200. IEEE, 2011.
- [48] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting on association for computational linguistics*, pages 311–318. Association for Computational Linguistics, 2002.
- [49] Max B Paulus, Dami Choi, Daniel Tarlow, Andreas Krause, and Chris J Maddison. Gradient estimation with stochastic softmax tricks. *arXiv preprint arXiv:2006.08063*, 2020.
- [50] Ratish Puduppully, Li Dong, and Mirella Lapata. Data-to-text generation with entity modeling. In *ACL*, 2019.
- [51] Yevgeniy Puzikov and Iryna Gurevych. E2e nlg challenge: Neural models vs. templates. In *Proceedings of the 11th International Conference on Natural Language Generation*, pages 463–471, 2018.
- [52] Uvais Qidwai and Mohamed Shakir. Ubiquitous arabic voice control device to assist people with disabilities. *2012 4th International Conference on Intelligent and Advanced Systems (ICIAS2012)*, 1:333–338, 2012.
- [53] Rajesh Ranganath, Sean Gerrish, and David M Blei. Black box variational inference. *arXiv preprint arXiv:1401.0118*, 2013.
- [54] Arsénio Reis, Dennis Paulino, Hugo Paredes, Isabel Barroso, Maria João Monteiro, Vitor Rodrigues, and João Barroso. Using intelligent personal assistants to assist the elderlies an evaluation of amazon alexa, google assistant, microsoft cortana, and apple siri. *2018 2nd International Conference on Technology and Innovation in Sports, Health and Wellbeing (TISHW)*, pages 1–5, 2018.
- [55] Danilo Jimenez Rezende, Shakir Mohamed, and Daan Wierstra. Stochastic backpropagation and approximate inference in deep generative models. *arXiv preprint arXiv:1401.4082*, 2014.
- [56] Tianxiao Shen, Tao Lei, Regina Barzilay, and Tommi Jaakkola. Style transfer from non-parallel text by cross-alignment. In *Advances in neural information processing systems*, pages 6830–6841, 2017.
- [57] David Silver, Aja Huang, Chris J Maddison, Arthur Guez, Laurent Sifre, George Van Den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, et al. Mastering the game of go with deep neural networks and tree search. *nature*, 529(7587):484, 2016.
- [58] Hong Sun and Ming Zhou. Joint learning of a dual SMT system for paraphrase generation. In *Proceedings of the 50th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers)*, pages 38–42, Jeju Island, Korea, July 2012. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/P12-2008>.
- [59] Charles Sutton, Andrew McCallum, et al. An introduction to conditional random fields. *Foundations and Trends® in Machine Learning*, 4(4):267–373, 2012.
- [60] George Tucker, Andriy Mnih, Chris J. Maddison, John Lawson, and Jascha Sohl-Dickstein. Rebar: Low-variance, unbiased gradient estimates for discrete latent variable models. *ArXiv*, abs/1703.07370, 2017.

- [61] Ramakrishna Vedantam, C Lawrence Zitnick, and Devi Parikh. Cider: Consensus-based image description evaluation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 4566–4575, 2015.
- [62] Sarah Wiegrefe and Yuval Pinter. Attention is not not explanation. In *EMNLP/IJCNLP*, 2019.
- [63] Ronald J Williams. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning*, 8(3-4):229–256, 1992.
- [64] Sam Wiseman, Stuart M Shieber, and Alexander M Rush. Learning neural templates for text generation. *arXiv preprint arXiv:1808.10122*, 2018.
- [65] Sang Michael Xie and Stefano Ermon. Differentiable subset sampling. *arXiv preprint arXiv:1901.10517*, 2019.
- [66] Pengcheng Yin, Chunting Zhou, Junxian He, and Graham Neubig. Structvae: Tree-structured latent variable models for semi-supervised semantic parsing. *arXiv preprint arXiv:1806.07832*, 2018.
- [67] Lantao Yu, Weinan Zhang, Jun Wang, and Yong Yu. Seqgan: Sequence generative adversarial nets with policy gradient. *ArXiv*, abs/1609.05473, 2017.
- [68] Jake Zhao, Yoon Kim, Kelly Zhang, Alexander M Rush, and Yann LeCun. Adversarially regularized autoencoders. *arXiv preprint arXiv:1706.04223*, 2017.